

winequality-red

February 24, 2024

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
```

```
[2]: df = pd.read_csv('winequality-red.csv')
```

```
[3]: print(df.head())
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
0	7.4	0.70	0.00	1.9	0.076	
1	7.8	0.88	0.00	2.6	0.098	
2	7.8	0.76	0.04	2.3	0.092	
3	11.2	0.28	0.56	1.9	0.075	
4	7.4	0.70	0.00	1.9	0.076	

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	\
0	11.0	34.0	0.9978	3.51	0.56	
1	25.0	67.0	0.9968	3.20	0.68	
2	15.0	54.0	0.9970	3.26	0.65	
3	17.0	60.0	0.9980	3.16	0.58	
4	11.0	34.0	0.9978	3.51	0.56	

	alcohol	quality
0	9.4	5
1	9.8	5
2	9.8	5
3	9.8	6
4	9.4	5

```
[4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
#   Column              Non-Null Count  Dtype
---  -
#   Column              Non-Null Count  Dtype
```

```

0    fixed acidity      1599 non-null    float64
1    volatile acidity   1599 non-null    float64
2    citric acid        1599 non-null    float64
3    residual sugar     1599 non-null    float64
4    chlorides          1599 non-null    float64
5    free sulfur dioxide 1599 non-null    float64
6    total sulfur dioxide 1599 non-null    float64
7    density            1599 non-null    float64
8    pH                 1599 non-null    float64
9    sulphates          1599 non-null    float64
10   alcohol            1599 non-null    float64
11   quality            1599 non-null    int64

```

dtypes: float64(11), int64(1)

memory usage: 150.0 KB

```
[5]: df.isnull().sum()
```

```

[5]: fixed acidity      0
     volatile acidity   0
     citric acid        0
     residual sugar     0
     chlorides          0
     free sulfur dioxide 0
     total sulfur dioxide 0
     density            0
     pH                 0
     sulphates          0
     alcohol            0
     quality            0
     dtype: int64

```

```
[6]: df.describe
```

```

[6]: <bound method NDFrame.describe of      fixed acidity  volatile acidity  citric
acid residual sugar  chlorides \
0                7.4          0.700      0.00      1.9      0.076
1                7.8          0.880      0.00      2.6      0.098
2                7.8          0.760      0.04      2.3      0.092
3               11.2          0.280      0.56      1.9      0.075
4                7.4          0.700      0.00      1.9      0.076
...          ...          ...          ...          ...          ...
1594            6.2          0.600      0.08      2.0      0.090
1595            5.9          0.550      0.10      2.2      0.062
1596            6.3          0.510      0.13      2.3      0.076
1597            5.9          0.645      0.12      2.0      0.075
1598            6.0          0.310      0.47      3.6      0.067

```

```

free sulfur dioxide  total sulfur dioxide  density  pH  sulphates \

```

0	11.0	34.0	0.99780	3.51	0.56
1	25.0	67.0	0.99680	3.20	0.68
2	15.0	54.0	0.99700	3.26	0.65
3	17.0	60.0	0.99800	3.16	0.58
4	11.0	34.0	0.99780	3.51	0.56
...
1594	32.0	44.0	0.99490	3.45	0.58
1595	39.0	51.0	0.99512	3.52	0.76
1596	29.0	40.0	0.99574	3.42	0.75
1597	32.0	44.0	0.99547	3.57	0.71
1598	18.0	42.0	0.99549	3.39	0.66

	alcohol	quality
0	9.4	5
1	9.8	5
2	9.8	5
3	9.8	6
4	9.4	5
...
1594	10.5	5
1595	11.2	6
1596	11.0	6
1597	10.2	5
1598	11.0	6

[1599 rows x 12 columns]>

```
[7]: df.describe()
```

```
[7]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar \
count	1599.000000	1599.000000	1599.000000	1599.000000
mean	8.319637	0.527821	0.270976	2.538806
std	1.741096	0.179060	0.194801	1.409928
min	4.600000	0.120000	0.000000	0.900000
25%	7.100000	0.390000	0.090000	1.900000
50%	7.900000	0.520000	0.260000	2.200000
75%	9.200000	0.640000	0.420000	2.600000
max	15.900000	1.580000	1.000000	15.500000

	chlorides	free sulfur dioxide	total sulfur dioxide	density \
count	1599.000000	1599.000000	1599.000000	1599.000000
mean	0.087467	15.874922	46.467792	0.996747
std	0.047065	10.460157	32.895324	0.001887
min	0.012000	1.000000	6.000000	0.990070
25%	0.070000	7.000000	22.000000	0.995600
50%	0.079000	14.000000	38.000000	0.996750
75%	0.090000	21.000000	62.000000	0.997835

max	0.611000	72.000000	289.000000	1.003690
-----	----------	-----------	------------	----------

	pH	sulphates	alcohol	quality
count	1599.000000	1599.000000	1599.000000	1599.000000
mean	3.311113	0.658149	10.422983	5.636023
std	0.154386	0.169507	1.065668	0.807569
min	2.740000	0.330000	8.400000	3.000000
25%	3.210000	0.550000	9.500000	5.000000
50%	3.310000	0.620000	10.200000	6.000000
75%	3.400000	0.730000	11.100000	6.000000
max	4.010000	2.000000	14.900000	8.000000

```
[8]: df.corr()
```

```
[8]:
```

	fixed acidity	volatile acidity	citric acid	\
fixed acidity	1.000000	-0.256131	0.671703	
volatile acidity	-0.256131	1.000000	-0.552496	
citric acid	0.671703	-0.552496	1.000000	
residual sugar	0.114777	0.001918	0.143577	
chlorides	0.093705	0.061298	0.203823	
free sulfur dioxide	-0.153794	-0.010504	-0.060978	
total sulfur dioxide	-0.113181	0.076470	0.035533	
density	0.668047	0.022026	0.364947	
pH	-0.682978	0.234937	-0.541904	
sulphates	0.183006	-0.260987	0.312770	
alcohol	-0.061668	-0.202288	0.109903	
quality	0.124052	-0.390558	0.226373	

	residual sugar	chlorides	free sulfur dioxide	\
fixed acidity	0.114777	0.093705	-0.153794	
volatile acidity	0.001918	0.061298	-0.010504	
citric acid	0.143577	0.203823	-0.060978	
residual sugar	1.000000	0.055610	0.187049	
chlorides	0.055610	1.000000	0.005562	
free sulfur dioxide	0.187049	0.005562	1.000000	
total sulfur dioxide	0.203028	0.047400	0.667666	
density	0.355283	0.200632	-0.021946	
pH	-0.085652	-0.265026	0.070377	
sulphates	0.005527	0.371260	0.051658	
alcohol	0.042075	-0.221141	-0.069408	
quality	0.013732	-0.128907	-0.050656	

	total sulfur dioxide	density	pH	sulphates	\
fixed acidity	-0.113181	0.668047	-0.682978	0.183006	
volatile acidity	0.076470	0.022026	0.234937	-0.260987	
citric acid	0.035533	0.364947	-0.541904	0.312770	
residual sugar	0.203028	0.355283	-0.085652	0.005527	


chlorides	0.047400	0.200632	-0.265026	0.371260
free sulfur dioxide	0.667666	-0.021946	0.070377	0.051658
total sulfur dioxide	1.000000	0.071269	-0.066495	0.042947
density	0.071269	1.000000	-0.341699	0.148506
pH	-0.066495	-0.341699	1.000000	-0.196648
sulphates	0.042947	0.148506	-0.196648	1.000000
alcohol	-0.205654	-0.496180	0.205633	0.093595
quality	-0.185100	-0.174919	-0.057731	0.251397

	alcohol	quality
fixed acidity	-0.061668	0.124052
volatile acidity	-0.202288	-0.390558
citric acid	0.109903	0.226373
residual sugar	0.042075	0.013732
chlorides	-0.221141	-0.128907
free sulfur dioxide	-0.069408	-0.050656
total sulfur dioxide	-0.205654	-0.185100
density	-0.496180	-0.174919
pH	0.205633	-0.057731
sulphates	0.093595	0.251397
alcohol	1.000000	0.476166
quality	0.476166	1.000000


```
[9]: corr_data = df.corr()
```

```
[10]: plt.figure(figsize = (9, 9))
sns.heatmap(corr_data, cbar = True, square= True, annot=True, fmt= '.2f',
            cmap='Blues')
plt.title('Heatmap')
```


```
[10]: Text(0.5, 1.0, 'Heatmap')
```



```
[11]: def hist_plots(df):
        plt.figure(figsize=(10,6))
        plt.hist(df)
        plt.title("Histogram")
        hist_plots(df['fixed acidity'])
```

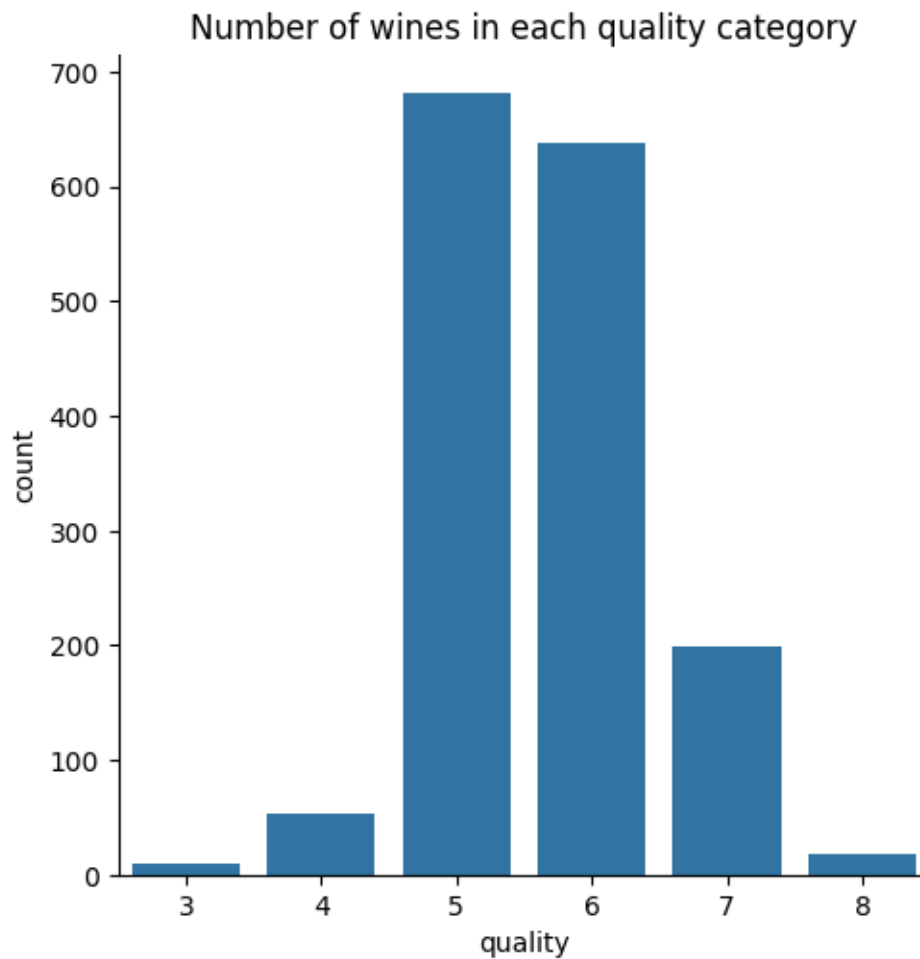


```
[12]: hist_plots(df['pH'])
```




```
[13]: # Number of wines in each quality category
sns.catplot(x='quality', data=df, kind='count')
plt.title('Number of wines in each quality category')
```

```
[13]: Text(0.5, 1.0, 'Number of wines in each quality category')
```




```
[14]: hist_plots(df['quality'])
```





```
[15]: # Barplot for quality vs volatile acidity  
sns.barplot(x = 'quality', y= 'volatile acidity', data = df)
```

```
[15]: <Axes: xlabel='quality', ylabel='volatile acidity'>
```




```
[16]: hist_plots(df['alcohol'])
```




```
[17]: #Barplot for quality vs alcohol  
sns.barplot(x = 'quality', y = 'alcohol', data = df)
```

```
[17]: <Axes: xlabel='quality', ylabel='alcohol'>
```




```
[18]: hist_plots(df['residual sugar'])
```




```
[19]: sns.kdeplot(df['volatile acidity'])
```

```
[19]: <Axes: xlabel='volatile acidity', ylabel='Density'>
```




```
[20]: sns.kdeplot(df['chlorides'])
```

```
[20]: <Axes: xlabel='chlorides', ylabel='Density'>
```




```
[21]: sns.kdeplot(df['density'])
```

```
[21]: <Axes: xlabel='density', ylabel='Density'>
```



```
[22]: df.hist(figsize=(11,11))  
plt.show()
```

```
[23]: X = df.drop('quality', axis= 1)
      y = df['quality'].apply(lambda y_value: 1 if y_value >= 7 else 0)
```

```
[24]: y.value_counts()
```

```
[24]: quality
0      1382
1       217
Name: count, dtype: int64
```

```
[25]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state= 3)
```

```
[26]: model = RandomForestClassifier(n_estimators=100, max_depth=5, random_state=1)
```

```
[27]: model.fit(X_train, y_train)
```

```
[27]: RandomForestClassifier(max_depth=5, random_state=1)
```

```
[28]: from sklearn.metrics import accuracy_score
```

```
[29]: X_test_preds = model.predict(X_test)  
test_accuracy = accuracy_score(y_test, X_test_preds)
```

```
[30]: print("Test Accuracy {:.2f}%".format(test_accuracy * 100))
```

```
Test Accuracy 91.56%
```